

Final_Code

May 13, 2025

```
[1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

EXPLORATORY DATA ANALYSIS BEGIN

```
[2]: #Apple stock dataset
file_apple = "apple_stock.csv"
apple_df = pd.read_csv(file_apple)
apple_df
```

```
[2]:
```

	Unnamed: 0	Adj Close	Close	High	Low	Open \
0	1980-12-12	0.098834	0.128348	0.128906	0.128348	0.128348
1	1980-12-15	0.093678	0.121652	0.122210	0.121652	0.122210
2	1980-12-16	0.086802	0.112723	0.113281	0.112723	0.113281
3	1980-12-17	0.088951	0.115513	0.116071	0.115513	0.115513
4	1980-12-18	0.091530	0.118862	0.119420	0.118862	0.118862
...
11102	2024-12-27	255.589996	255.589996	258.700012	253.059998	257.829987
11103	2024-12-30	252.199997	252.199997	253.500000	250.750000	252.229996
11104	2024-12-31	250.419998	250.419998	253.279999	249.429993	252.440002
11105	2025-01-02	243.850006	243.850006	249.100006	241.820007	248.929993
11106	2025-01-03	243.860001	243.860001	244.179993	241.889999	243.369995
	Volume					
0	469033600					
1	175884800					
2	105728000					
3	86441600					
4	73449600					
...	...					
11102	42355300					
11103	35557500					
11104	39480700					
11105	55558000					
11106	15135053					

[11107 rows x 7 columns]

```
[3]: #Microsoft stock dataset
file_msft = "MSFT_stock.csv"
msft_df = pd.read_csv(file_msft)
msft_df
```

```
[3]:
```

	Date	Open	High	Low	Close	Adj Close \
0	1986-03-13	0.088542	0.101563	0.088542	0.097222	0.059827
1	1986-03-14	0.097222	0.102431	0.097222	0.100694	0.061963
2	1986-03-17	0.100694	0.103299	0.100694	0.102431	0.063032
3	1986-03-18	0.102431	0.103299	0.098958	0.099826	0.061429
4	1986-03-19	0.099826	0.100694	0.097222	0.098090	0.060361
...
9795	2025-01-28	434.600006	448.380005	431.380005	447.200012	447.200012
9796	2025-01-29	446.690002	446.880005	440.399994	442.329987	442.329987
9797	2025-01-30	418.769989	422.859985	413.160004	414.989990	414.989990
9798	2025-01-31	418.980011	420.690002	414.910004	415.059998	415.059998
9799	2025-02-03	411.600006	415.410004	408.660004	410.920013	410.920013
	Volume					
0	1031788800					
1	308160000					
2	133171200					
3	67766400					
4	47894400					
...	...					
9795	23491700					
9796	23581400					
9797	54586300					
9798	34223400					
9799	25580600					

[9800 rows x 7 columns]

```
[4]: #rename date columnn make sure it is Datetime formayt
apple_df.rename(columns={"Unnamed: 0": "date"}, inplace=True)
apple_df['date'] = pd.to_datetime(apple_df['date'])
print("Apple Table: \n\n", apple_df.head())

msft_df.rename(columns={"Date": "date"}, inplace=True)
msft_df['date'] = pd.to_datetime(msft_df['date'])
print("\n\nMicrosoft Table: \n\n", msft_df.head())
```

Apple Table:

	date	Adj Close	Close	High	Low	Open	Volume
0	1980-12-12	0.098834	0.128348	0.128906	0.128348	0.128348	469033600

1	1980-12-15	0.093678	0.121652	0.122210	0.121652	0.122210	175884800
2	1980-12-16	0.086802	0.112723	0.113281	0.112723	0.113281	105728000
3	1980-12-17	0.088951	0.115513	0.116071	0.115513	0.115513	86441600
4	1980-12-18	0.091530	0.118862	0.119420	0.118862	0.118862	73449600

Microsoft Table:

	date	Open	High	Low	Close	Adj Close	Volume
0	1986-03-13	0.088542	0.101563	0.088542	0.097222	0.059827	1031788800
1	1986-03-14	0.097222	0.102431	0.097222	0.100694	0.061963	308160000
2	1986-03-17	0.100694	0.103299	0.100694	0.102431	0.063032	133171200
3	1986-03-18	0.102431	0.103299	0.098958	0.099826	0.061429	67766400
4	1986-03-19	0.099826	0.100694	0.097222	0.098090	0.060361	47894400

```
[5]: #creating new columns daily return (how much adj_close % changes) and
      ↪Volatility (5 day/one trading week std of daily return)
apple_df['Daily Return'] = apple_df['Adj Close'].pct_change()*100
msft_df['Daily Return'] = msft_df['Adj Close'].pct_change()*100

apple_df['Volatility_5'] = apple_df['Daily Return'].rolling(window=5).std()
msft_df['Volatility_5'] = msft_df['Daily Return'].rolling(window=5).std()

print("Apple Table: \n\n", apple_df.tail())
print("\n\nMicrosoft Table: \n\n", msft_df.tail())
```

Apple Table:

	date	Adj Close	Close	High	Low	Open \
11102	2024-12-27	255.589996	255.589996	258.700012	253.059998	257.829987
11103	2024-12-30	252.199997	252.199997	253.500000	250.750000	252.229996
11104	2024-12-31	250.419998	250.419998	253.279999	249.429993	252.440002
11105	2025-01-02	243.850006	243.850006	249.100006	241.820007	248.929993
11106	2025-01-03	243.860001	243.860001	244.179993	241.889999	243.369995

	Volume	Daily Return	Volatility_5
11102	42355300	-1.324219	1.195941
11103	35557500	-1.326343	1.103469
11104	39480700	-0.705789	1.085592
11105	55558000	-2.623589	1.070336
11106	15135053	0.004099	0.968502

Microsoft Table:

	date	Open	High	Low	Close	Adj Close \
9795	2025-01-28	434.600006	448.380005	431.380005	447.200012	447.200012
9796	2025-01-29	446.690002	446.880005	440.399994	442.329987	442.329987

9797	2025-01-30	418.769989	422.859985	413.160004	414.989990	414.989990
9798	2025-01-31	418.980011	420.690002	414.910004	415.059998	415.059998
9799	2025-02-03	411.600006	415.410004	408.660004	410.920013	410.920013

	Volume	Daily Return	Volatility_5
9795	23491700	2.908693	2.576610
9796	23581400	-1.089004	1.901359
9797	54586300	-6.180905	3.267812
9798	34223400	0.016870	3.317358
9799	25580600	-0.997442	3.283986

```
[6]: #Statistical breakdown
print("Apple Summary Stats: \n\n", apple_df[['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume', 'Daily Return', 'Volatility_5']].describe())
print("\n\nMicrosoft Summary Stats: \n\n", msft_df[['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume', 'Daily Return', 'Volatility_5']].describe())
```

Apple Summary Stats:

	Open	High	Low	Close	Adj Close \
count	11107.000000	11107.000000	11107.000000	11107.000000	11107.000000
mean	24.339076	24.598169	24.092608	24.357607	23.522229
std	50.166818	50.691902	49.682631	50.217498	49.767881
min	0.049665	0.049665	0.049107	0.049107	0.037815
25%	0.300090	0.306362	0.292411	0.300290	0.243402
50%	0.542679	0.553393	0.534598	0.542411	0.446682
75%	21.367679	21.569285	21.115715	21.397143	18.260086
max	258.190002	260.100006	257.630005	259.019989	259.019989

	Volume	Daily Return	Volatility_5
count	1.110700e+04	11106.000000	11102.000000
mean	3.154341e+08	0.109604	2.304649
std	3.348735e+08	2.778277	1.562694
min	0.000000e+00	-51.869200	0.075523
25%	1.111164e+08	-1.260233	1.266991
50%	2.036944e+08	0.000000	1.962168
75%	3.960418e+08	1.437862	2.946773
max	7.421641e+09	33.228009	24.495311

Microsoft Summary Stats:

	Open	High	Low	Close	Adj Close \
count	9800.000000	9800.000000	9800.000000	9800.000000	9800.000000
mean	63.173728	63.812493	62.513818	63.187157	57.451679
std	98.682750	99.575073	97.725918	98.697881	98.714771
min	0.088542	0.092014	0.088542	0.090278	0.055554
25%	5.898438	5.976563	5.794922	5.880860	3.618860

50%	27.436250	27.770000	27.200001	27.490000	19.266710
75%	47.597813	48.145000	47.062500	47.592500	40.066001
max	467.000000	468.350006	464.459991	467.559998	465.786438

	Volume	Daily Return	Volatility_5
count	9.800000e+03	9799.000000	9795.000000
mean	5.630897e+07	0.112386	1.756862
std	3.812127e+07	2.102810	1.206792
min	2.304000e+06	-30.115887	0.068665
25%	3.141062e+07	-0.910068	0.964287
50%	4.946235e+07	0.037737	1.479103
75%	7.027870e+07	1.119693	2.207315
max	1.031789e+09	19.565212	18.018816

```
[7]: #Notice NaNs
apple_df.head()
```

```
[7]:      date  Adj Close  Close  High  Low  Open  Volume \
0 1980-12-12  0.098834  0.128348  0.128906  0.128348  0.128348  469033600
1 1980-12-15  0.093678  0.121652  0.122210  0.121652  0.122210  175884800
2 1980-12-16  0.086802  0.112723  0.113281  0.112723  0.113281  105728000
3 1980-12-17  0.088951  0.115513  0.116071  0.115513  0.115513   86441600
4 1980-12-18  0.091530  0.118862  0.119420  0.118862  0.118862   73449600
```

	Daily Return	Volatility_5
0	NaN	NaN
1	-5.217063	NaN
2	-7.339813	NaN
3	2.475121	NaN
4	2.899232	NaN

```
[8]: #Notice NaNs
msft_df.head()
```

```
[8]:      date  Open  High  Low  Close  Adj Close  Volume \
0 1986-03-13  0.088542  0.101563  0.088542  0.097222  0.059827  1031788800
1 1986-03-14  0.097222  0.102431  0.097222  0.100694  0.061963  308160000
2 1986-03-17  0.100694  0.103299  0.100694  0.102431  0.063032  133171200
3 1986-03-18  0.102431  0.103299  0.098958  0.099826  0.061429   67766400
4 1986-03-19  0.099826  0.100694  0.097222  0.098090  0.060361   47894400
```

	Daily Return	Volatility_5
0	NaN	NaN
1	3.571207	NaN
2	1.725111	NaN
3	-2.543234	NaN
4	-1.738989	NaN

```
[9]: #only Nans are in the first five rows cause it uses previous 5 days before that
      ↪ day to calculate, so drop those
apple_df.dropna(inplace=True)
msft_df.dropna(inplace=True)
```

```
[10]: apple_df.head()
```

```
[10]:
```

	date	Adj Close	Close	High	Low	Open	Volume	\
5	1980-12-19	0.097116	0.126116	0.126674	0.126116	0.126116	48630400	
6	1980-12-22	0.101842	0.132254	0.132813	0.132254	0.132254	37363200	
7	1980-12-23	0.106140	0.137835	0.138393	0.137835	0.137835	46950400	
8	1980-12-24	0.111726	0.145089	0.145647	0.145089	0.145089	48003200	
9	1980-12-26	0.122039	0.158482	0.159040	0.158482	0.158482	55574400	

	Daily Return	Volatility_5
5	6.102858	5.758344
6	4.867013	5.317928
7	4.219907	1.474397
8	5.262747	1.202232
9	9.230927	1.963757

```
[11]: msft_df.head()
```

```
[11]:
```

	date	Open	High	Low	Close	Adj Close	Volume	\
5	1986-03-20	0.098090	0.098090	0.094618	0.095486	0.058758	58435200	
6	1986-03-21	0.095486	0.097222	0.091146	0.092882	0.057156	59990400	
7	1986-03-24	0.092882	0.092882	0.089410	0.090278	0.055554	65289600	
8	1986-03-25	0.090278	0.092014	0.089410	0.092014	0.056622	32083200	
9	1986-03-26	0.092014	0.095486	0.091146	0.094618	0.058224	22752000	

	Daily Return	Volatility_5
5	-2.654749	2.816513
6	-2.727116	1.893905
7	-2.803508	0.432563
8	1.922941	2.016293
9	2.830012	2.814910

```
[12]: #breakdown of cleaned up datasets now
print("Apple description:")
print(apple_df.describe())
print("\nApple data types:")
print(apple_df.dtypes)
print("\nApple missing values:")
print(apple_df.isna().sum())

print("\nMicrosoft description:")
print(msft_df.describe())
```

```
print("\nMicrosoft data types:")
print(msft_df.dtypes)
print("\nMicrosoft missing values:")
print(msft_df.isna().sum())
```

Apple description:

	date	Adj Close	Close \
count	11102	11102.000000	11102.000000
mean	2002-12-18 06:22:14.678436352	23.532781	24.368523
min	1980-12-19 00:00:00	0.037815	0.049107
25%	1991-12-11 06:00:00	0.243700	0.300781
50%	2002-12-12 12:00:00	0.447646	0.542411
75%	2013-12-22 06:00:00	18.272139	21.406875
max	2025-01-03 00:00:00	259.019989	259.019989
std	NaN	49.776603	50.226171

	High	Low	Open	Volume	Daily Return \
count	11102.000000	11102.000000	11102.000000	1.110200e+04	11102.000000
mean	24.609193	24.103404	24.349983	3.154942e+08	0.110290
min	0.049665	0.049107	0.049665	0.000000e+00	-51.869200
25%	0.306920	0.293292	0.300223	1.111544e+08	-1.259590
50%	0.553571	0.535435	0.544464	2.037560e+08	0.000000
75%	21.588483	21.131250	21.400893	3.960593e+08	1.437291
max	260.100006	257.630005	258.190002	7.421641e+09	33.228009
std	50.700654	49.691213	50.175480	3.349223e+08	2.777201

	Volatility_5
count	11102.000000
mean	2.304649
min	0.075523
25%	1.266991
50%	1.962168
75%	2.946773
max	24.495311
std	1.562694

Apple data types:

```
date          datetime64[ns]
Adj Close      float64
Close          float64
High          float64
Low           float64
Open          float64
Volume        int64
Daily Return   float64
Volatility_5   float64
dtype: object
```

Apple missing values:

```
date          0
Adj Close     0
Close         0
High          0
Low           0
Open          0
Volume        0
Daily Return  0
Volatility_5  0
dtype: int64
```

Microsoft description:

	date	Open	High	Low \
count	9795	9795.000000	9795.000000	9795.000000
mean	2005-08-17 20:16:23.522205056	63.205926	63.845015	62.545680
min	1986-03-20 00:00:00	0.090278	0.092014	0.089410
25%	1995-11-23 00:00:00	5.921875	6.015625	5.812500
50%	2005-08-16 00:00:00	27.450001	27.775000	27.200001
75%	2015-05-09 12:00:00	47.617500	48.173750	47.076250
max	2025-02-03 00:00:00	467.000000	468.350006	464.459991
std	NaN	98.697642	99.590079	97.740681

	Close	Adj Close	Volume	Daily Return	Volatility_5
count	9795.000000	9795.000000	9.795000e+03	9795.000000	9795.000000
mean	63.219361	57.480975	5.617551e+07	0.112329	1.756862
min	0.090278	0.055554	2.304000e+06	-30.115887	0.068665
25%	5.906250	3.634482	3.140820e+07	-0.909028	0.964287
50%	27.500000	19.271568	4.945560e+07	0.037737	1.479103
75%	47.605000	40.090405	7.025070e+07	1.119149	2.207315
max	467.559998	465.786438	7.886880e+08	19.565212	18.018816
std	98.712775	98.731446	3.673823e+07	2.102631	1.206792

Microsoft data types:

```
date          datetime64[ns]
Open          float64
High          float64
Low           float64
Close         float64
Adj Close     float64
Volume        int64
Daily Return  float64
Volatility_5  float64
dtype: object
```

Microsoft missing values:

```
date          0
Open          0
```



```

High          0
Low           0
Close         0
Adj Close     0
Volume        0
Daily Return  0
Volatility_5  0
dtype: int64


```

```

[13]: #Visualization - Histograms of daily returns
plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.histplot(apple_df['Daily Return'], color='blue')
plt.title('Apple Daily Returns Distribution')
plt.xlabel('Daily Return(%)')

plt.subplot(1, 2, 2)
sns.histplot(msft_df['Daily Return'], color='green')
plt.title('Microsoft Daily Returns Distribution')
plt.xlabel('Daily Return (%)')
plt.tight_layout()
plt.show()

```




```

[14]: #Visualization - boxplots of volume traded
plt.figure(figsize=(15, 5))
plt.subplot(1, 2, 1)
sns.boxplot(x=apple_df['Volume'], color='skyblue')
plt.title('Apple Volume Distribution')


plt.subplot(1, 2, 2)
sns.boxplot(x=msft_df['Volume'], color='lightgreen')
plt.title('Microsoft Volume Distribution')
plt.tight_layout()
plt.show()

```




```
[15]: #Visualization - line charts of adj closing price over time
plt.figure(figsize=(15, 5))
plt.plot(apple_df['date'], apple_df['Adj Close'], label='Apple', color='blue',
         ↪linewidth=2)
plt.plot(msft_df['date'], msft_df['Adj Close'], label='Microsoft',
         ↪color='green', linewidth=2)

plt.title('Adjusted Close Price Over Time')
plt.xlabel('date')
plt.ylabel('Adjusted Close Price')
plt.legend()
plt.show()
```



```
[16]: #Visualization - Volatility of daily returns over a 5 day period (1 trading
         ↪week)
plt.plot(apple_df['date'], apple_df['Volatility_5'], label='Apple 5-Day
         ↪Volatility', color='blue')
```

```
plt.plot(msft_df['date'], msft_df['Volatility_5'], label='Microsoft 5-Day Volatility', color='green')
plt.title('5-Day Rolling Volatility of Daily Returns')
plt.xlabel('date')
plt.ylabel('Volatility')
plt.legend()
plt.show()
```



EXPLORATORY DATA ANALYSIS END

Logistic regression model START: First have to choose events that we want to look at and put them into two separate "event types".

1 - technological events 2 - Government policy events

apple tech events - announcement of iPhone, announcement of iPad, announcement of M1 chip
apple policy events - US-China tariffs, Trump's first time in office, US-EU antitrust probe

microsoft tech events - release of Windows 95, announcement of Azure, announcement of Xbox
microsoft policy events - antitrust lawsuit filed, antitrust lawsuit settled

```
[17]: #Mark events in dataset first: 0 = normal day, 1 = post-technological innovation event, 2 = post-government policy event
```

```

apple_events = {
    '2007-01-09': 1,    #iPhone
    '2010-01-27': 1,    #iPad
    '2020-11-10': 1,    #M1 chip
    '2018-03-22': 2,    #US-China tariffs
    '2020-06-16': 2     #US-EU antitrust probe
}

msft_events = {
    '1995-08-24': 1,    # Windows 95
    '2010-02-01': 1,    # Azure
    '2000-03-10': 1,    # Xbox
    '1998-05-18': 2,    # US antitrust lawsuit filed
    '2001-11-02': 2     # Final antitrust settled
}

def mark_event_window(df, event_dates, column='date', window=5):
    df['event_type'] = 0    #every normal day is 0

    for event_date_str, event_code in event_dates.items():
        event_date = pd.to_datetime(event_date_str)

        #first day after the event, event may have been on weekend so next
        ↪available "trading day"
        future_dates = df[column][df[column] >= event_date]
        if len(future_dates) == 0:
            continue

        start_idx = future_dates.index[0]
        end_idx = start_idx + window

        #if event is near end of df
        if end_idx < len(df):
            df.loc[start_idx:end_idx, 'event_type'] = event_code

    return df

apple_df = mark_event_window(apple_df, apple_events)
msft_df = mark_event_window(msft_df, msft_events)

```

```

[18]: #want to make sure the events dates were correctly edited in each dataset
      #Apple
      print("Apple Event Type Count:")
      print(apple_df['event_type'].value_counts())
      print("Apple - Technology Events (1)")
      display(apple_df[apple_df['event_type'] == 1][['date', 'event_type']])

```

```
print("Apple - Government Policy Events (2)")
display(apple_df[apple_df['event_type'] == 2][['date', 'event_type']])
```

Apple Event Type Count:

event_type

0 11072

1 18

2 12

Name: count, dtype: int64

Apple - Technology Events (1)

	date	event_type
6579	2007-01-09	1
6580	2007-01-10	1
6581	2007-01-11	1
6582	2007-01-12	1
6583	2007-01-16	1
6584	2007-01-17	1
7347	2010-01-27	1
7348	2010-01-28	1
7349	2010-01-29	1
7350	2010-02-01	1
7351	2010-02-02	1
7352	2010-02-03	1
10064	2020-11-10	1
10065	2020-11-11	1
10066	2020-11-12	1
10067	2020-11-13	1
10068	2020-11-16	1
10069	2020-11-17	1

Apple - Government Policy Events (2)

	date	event_type
9399	2018-03-22	2
9400	2018-03-23	2
9401	2018-03-26	2
9402	2018-03-27	2
9403	2018-03-28	2
9404	2018-03-29	2
9961	2020-06-16	2
9962	2020-06-17	2
9963	2020-06-18	2
9964	2020-06-19	2
9965	2020-06-22	2
9966	2020-06-23	2

[19]: apple_df

```
[19]:
```

	date	Adj Close	Close	High	Low	Open \
5	1980-12-19	0.097116	0.126116	0.126674	0.126116	0.126116
6	1980-12-22	0.101842	0.132254	0.132813	0.132254	0.132254
7	1980-12-23	0.106140	0.137835	0.138393	0.137835	0.137835
8	1980-12-24	0.111726	0.145089	0.145647	0.145089	0.145089
9	1980-12-26	0.122039	0.158482	0.159040	0.158482	0.158482
...
11102	2024-12-27	255.589996	255.589996	258.700012	253.059998	257.829987
11103	2024-12-30	252.199997	252.199997	253.500000	250.750000	252.229996
11104	2024-12-31	250.419998	250.419998	253.279999	249.429993	252.440002
11105	2025-01-02	243.850006	243.850006	249.100006	241.820007	248.929993
11106	2025-01-03	243.860001	243.860001	244.179993	241.889999	243.369995

	Volume	Daily Return	Volatility_5	event_type
5	48630400	6.102858	5.758344	0
6	37363200	4.867013	5.317928	0
7	46950400	4.219907	1.474397	0
8	48003200	5.262747	1.202232	0
9	55574400	9.230927	1.963757	0
...
11102	42355300	-1.324219	1.195941	0
11103	35557500	-1.326343	1.103469	0
11104	39480700	-0.705789	1.085592	0
11105	55558000	-2.623589	1.070336	0
11106	15135053	0.004099	0.968502	0

[11102 rows x 10 columns]

```
[20]: #Microsoft
print("\nMicrosoft Event Type Counts:")
print(msft_df['event_type'].value_counts())
print("Microsoft - Technology Events (1)")
display(msft_df[msft_df['event_type'] == 1][['date', 'event_type']])
print("Microsoft - Government Policy Events (2)")
display(msft_df[msft_df['event_type'] == 2][['date', 'event_type']])
```

Microsoft Event Type Counts:

event_type

0 9765

1 18

2 12

Name: count, dtype: int64

Microsoft - Technology Events (1)

	date	event_type
--	------	------------

2390	1995-08-24	1
------	------------	---

2391	1995-08-25	1
------	------------	---

2392	1995-08-28	1
2393	1995-08-29	1
2394	1995-08-30	1
2395	1995-08-31	1
3537	2000-03-10	1
3538	2000-03-13	1
3539	2000-03-14	1
3540	2000-03-15	1
3541	2000-03-16	1
3542	2000-03-17	1
6024	2010-02-01	1
6025	2010-02-02	1
6026	2010-02-03	1
6027	2010-02-04	1
6028	2010-02-05	1
6029	2010-02-08	1

Microsoft - Government Policy Events (2)

	date	event_type
3079	1998-05-18	2
3080	1998-05-19	2
3081	1998-05-20	2
3082	1998-05-21	2
3083	1998-05-22	2
3084	1998-05-26	2
3950	2001-11-02	2
3951	2001-11-05	2
3952	2001-11-06	2
3953	2001-11-07	2
3954	2001-11-08	2
3955	2001-11-09	2

msft_df

```
[21]: #functions for logistic regression
def sigmoid(z):
    z = np.clip(z, -500, 500) # prevents overflow
    return 1 / (1 + np.exp(-z))

#train model utilizng gradient descent
def fit_logistic(X, y, lr=0.01, n_iter=10000):
    weights = np.zeros(X.shape[1])
    bias = 0

    for _ in range(n_iter):
        linear = np.dot(X, weights) + bias
        probs = sigmoid(linear)
```

```

    error = probs - y
    dweigh = np.dot(X.T, error) / len(y)
    dbias = np.mean(error)

    weights -= lr * dweigh
    bias -= lr * dbias

    return weights, bias

#predict prob of stock daily return being positive
def predict_proba(X, weights, bias):
    return sigmoid(np.dot(X, weights) + bias)

#predict up or down based off of predict_proba. >= 0.5 is up <= 0.5 is down on
↳ the day
def predict(X, weights, bias, threshold=0.5):
    return (predict_proba(X, weights, bias) >= threshold).astype(int)

```

```

[22]: def train_model(df):
    train_data = df[df['event_type'] == 0]
    X_train = train_data[features].values
    y_train = train_data['Target'].values

    weights, bias = fit_logistic(X_train, y_train)
    return weights, bias

#make predictions during blocked out time period for event types
def evaluate_model(df, weights, bias, event_type):
    test_data = df[df['event_type'] == event_type]
    X_test = test_data[features].values
    y_test = test_data['Target'].values

    if len(y_test) == 0:
        return None, 0

    y_pred = predict(X_test, weights, bias)
    accuracy = np.mean(y_pred == y_test)
    return accuracy

```

```

[23]: features = ['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume', 'Daily_
↳ Return', 'Volatility_5']

#target, if = 1 then daily return is positive and vice versa for 0
apple_df['Target'] = (apple_df['Daily Return'].shift(-1) > 0).astype(int)
msft_df['Target'] = (msft_df['Daily Return'].shift(-1) > 0).astype(int)

#model is trained on normal days when event_typ = 0

```



```

apple_weights, apple_bias = train_model(apple_df)
msft_weights, msft_bias = train_model(msft_df)

#then makes predictions based off the time period we blocked out for each event
↪type
apple_tech_acc = evaluate_model(apple_df, apple_weights, apple_bias, 1)
apple_policy_acc = evaluate_model(apple_df, apple_weights, apple_bias, 2)

msft_tech_acc = evaluate_model(msft_df, msft_weights, msft_bias, 1)
msft_policy_acc = evaluate_model(msft_df, msft_weights, msft_bias, 2)

#accuracy results for each event type
print("Apple Tech Events - Accuracy:", round(apple_tech_acc, 3))
print("Apple Policy Events - Accuracy:", round(apple_policy_acc, 3))

print("Microsoft Tech Events - Accuracy:", round(msft_tech_acc, 3))
print("Microsoft Policy Events - Accuracy:", round(msft_policy_acc, 3))

```

Apple Tech Events - Accuracy: 0.444
 Apple Policy Events - Accuracy: 0.417
 Microsoft Tech Events - Accuracy: 0.444
 Microsoft Policy Events - Accuracy: 0.667

[24]:

```

#EXTRA VERIFICATION FOR PAPER
from sklearn.metrics import accuracy_score, precision_score, recall_score,
    ↪f1_score, confusion_matrix

#Apple Tech events
apple_tech_data = apple_df[apple_df['event_type'] == 1]
X_apple_tech = apple_tech_data[features].values
y_apple_tech = apple_tech_data['Target'].values
y_pred_apple_tech = predict(X_apple_tech, apple_weights, apple_bias)

print("\nApple Tech Event Verification:")
print("Precision:", round(precision_score(y_apple_tech, y_pred_apple_tech), 3))
print("F1 Score:", round(f1_score(y_apple_tech, y_pred_apple_tech), 3))
print("Confusion Matrix:\n", confusion_matrix(y_apple_tech, y_pred_apple_tech))

#Apple Policy Events
apple_policy_data = apple_df[apple_df['event_type'] == 2]
X_apple_policy = apple_policy_data[features].values
y_apple_policy = apple_policy_data['Target'].values
y_pred_apple_policy = predict(X_apple_policy, apple_weights, apple_bias)

print("\nApple Policy Event Verification:")
print("Precision:", round(precision_score(y_apple_policy, y_pred_apple_policy),
    ↪3))

```

```

print("F1 Score:", round(f1_score(y_apple_policy, y_pred_apple_policy), 3))
print("Confusion Matrix:\n", confusion_matrix(y_apple_policy,
↪y_pred_apple_policy))

#Microsoft Tech Events
msft_tech_data = msft_df[msft_df['event_type'] == 1]
X_msft_tech = msft_tech_data[features].values
y_msft_tech = msft_tech_data['Target'].values
y_pred_msft_tech = predict(X_msft_tech, msft_weights, msft_bias)

print("\nMicrosoft Tech Event Verification:")
print("Precision:", round(precision_score(y_msft_tech, y_pred_msft_tech), 3))
print("F1 Score:", round(f1_score(y_msft_tech, y_pred_msft_tech), 3))
print("Confusion Matrix:\n", confusion_matrix(y_msft_tech, y_pred_msft_tech))

# Microsoft Policy Events
msft_policy_data = msft_df[msft_df['event_type'] == 2]
X_msft_policy = msft_policy_data[features].values
y_msft_policy = msft_policy_data['Target'].values
y_pred_msft_policy = predict(X_msft_policy, msft_weights, msft_bias)

print("\nMicrosoft Policy Event Verification:")
print("Precision:", round(precision_score(y_msft_policy, y_pred_msft_policy),
↪3))
print("F1 Score:", round(f1_score(y_msft_policy, y_pred_msft_policy), 3))
print("Confusion Matrix:\n", confusion_matrix(y_msft_policy,
↪y_pred_msft_policy))

```

Apple Tech Event Verification:

Precision: 0.444

F1 Score: 0.615

Confusion Matrix:

```
[[ 0 10]
 [ 0  8]]
```

Apple Policy Event Verification:

Precision: 0.417

F1 Score: 0.588

Confusion Matrix:

```
[[0 7]
 [0 5]]
```

Microsoft Tech Event Verification:

Precision: 0.444

F1 Score: 0.615

Confusion Matrix:

```
[[ 0 10]
 [ 0  8]]
```

Microsoft Policy Event Verification:

Precision: 0.667

F1 Score: 0.8

Confusion Matrix:

```
[[0 4]
 [0 8]]
```

[]: